Notebook 2 — Intelligence Algorithms

Predictive Logistics Engine — Mumbai Fleet Simulation

What this notebook demonstrates:

1. **Kalman Filter** — GPS jitter detection (Signal 1 of 3-signal ensemble)
2. **Isolation Forest** — Sensor anomaly detection (Van #402 incident)
3. **XGBoost ETA Model** — Travel time prediction with Friday regime detection
4. **H3 Geospatial Indexing** — Peer-van corroboration logic
5. **3-Signal Ensemble** — Final jitter vs traffic stall classification

Run Notebook 1 first to generate the data files.

```
!pip install pandas numpy matplotlib seaborn scikit-learn xgboost scipy --quiet
```

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import seaborn as sns
from sklearn.ensemble import IsolationForest
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, r2_score
from scipy.ndimage import uniform_filter1d
import xgboost as xgb
import warnings
warnings.filterwarnings('ignore')

plt.rcParams.update({
    'figure.facecolor':'#0A1628','axes.facecolor':'#132040',
    'axes.edgecolor':'#1E3A5F','axes.labelcolor':'#94A3B8',
    'xtick.color':'#94A3B8','ytick.color':'#94A3B8',
    'text.color':'#FFFFFF','grid.color':'#1E3A5F','grid.alpha':0.4,
})

gps_df    = pd.read_csv('sample_data/gps_telemetry.csv')
sensor_df = pd.read_csv('sample_data/sensor_telemetry.csv')

gps_df['timestamp_dt']    = pd.to_datetime(gps_df['timestamp_utc'], unit='ms')
sensor_df['timestamp_dt'] = pd.to_datetime(sensor_df['timestamp_utc'], unit='ms
```

```
print(f"GPS rows: {len(gps_df):,}  |  Sensor rows: {len(sensor_df):,}")
print("Data loaded.")
```

```
GPS rows: 117,600  |  Sensor rows: 117,600
Data loaded.
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Algorithm 1 — Kalman Filter: GPS Jitter Detection

The Kalman filter is a linear quadratic estimator. It fuses noisy GPS observations with a vehicle motion model (speed + heading) to produce a smoothed position with an uncertainty covariance.

Key insight: When uncertainty covariance spikes → likely GPS noise. When covariance is low AND speed near-zero → genuine traffic stall.

```
def kalman_smooth_positions(lats, lons, speeds, headings, accuracies):
    """
    Simplified 2D Kalman filter for position smoothing.
    State: [lat, lon]  Observation: raw GPS position
    Process noise tuned to van dynamics (~5km/h uncertainty per step)
    """
    n = len(lats)
    smoothed_lats = np.zeros(n)
    smoothed_lons  = np.zeros(n)
    covariances    = np.zeros(n)

    # Init
    x = np.array([lats[0], lons[0]])
    P = np.eye(2) * 0.001  # initial covariance

    # Noise matrices
    Q = np.eye(2) * 1e-7   # process noise (van movement uncertainty)

    for i in range(n):
        # Observation noise based on gps_accuracy_m
        r_val  = (accuracies[i] / 111000) ** 2
        R      = np.eye(2) * r_val

        # Predict
        P = P + Q

        # Update (Kalman gain)
        K = P @ np.linalg.inv(P + R)
        z = np.array([lats[i], lons[i]])
```

```
        x = x + K @ (z - x)
        P = (np.eye(2) - K) @ P

        smoothed_lats[i] = x[0]
        smoothed_lons[i]  = x[1]
        covariances[i]    = np.trace(P)  # total uncertainty

    return smoothed_lats, smoothed_lons, covariances

# Apply to one jitter van and one normal van
jitter_van = gps_df[gps_df['is_jitter_van']==True]['van_id'].iloc[0]
normal_van = gps_df[gps_df['is_jitter_van']==False]['van_id'].iloc[0]

for label, van_id in [('jitter', jitter_van), ('normal', normal_van)]:
    v = gps_df[gps_df['van_id']==van_id].sort_values('timestamp_utc').head(100)
    sl, slon, cov = kalman_smooth_positions(
        v['lat'].values, v['lon'].values,
        v['speed_kmh'].values, v['heading_deg'].values, v['gps_accuracy_m'].val
    )
    gps_df.loc[v.index, f'smooth_lat_{label}'] = sl
    gps_df.loc[v.index, f'smooth_lon_{label}'] = slon
    gps_df.loc[v.index, 'kalman_covariance']   = cov

print(f"Kalman filter applied.")
print(f"Jitter van: {jitter_van}  |  Normal van: {normal_van}")
```

```
Kalman filter applied.
Jitter van: VAN-0008  |  Normal van: VAN-0001
```

```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

for ax, (label, van_id, is_j) in zip(axes, [('Jitter van (faulty GPS)', jitter_
                                             ('Normal van (clean GPS)', normal
    v = gps_df[gps_df['van_id']==van_id].sort_values('timestamp_utc').head(100)
    sl, slon, cov = kalman_smooth_positions(
        v['lat'].values, v['lon'].values,
        v['speed_kmh'].values, v['heading_deg'].values, v['gps_accuracy_m'].val
    )
    ax.scatter(v['lon'], v['lat'], s=8, alpha=0.5,
               color='#EF4444' if is_j else '#94A3B8',
               label='Raw GPS positions', zorder=2)
    ax.plot(slon, sl, color='#0EA5E9', linewidth=2.5,
            label='Kalman smoothed path', zorder=3)
    ax.set_title(label, fontsize=12, pad=10)
    ax.set_xlabel('Longitude'); ax.set_ylabel('Latitude')
    ax.legend(fontsize=9)
    ax.grid(True, alpha=0.3)

    noise_str = f"GPS accuracy: {v['gps_accuracy_m'].mean():.0f}m avg" if is_j
    ax.text(0.02, 0.05, noise_str, transform=ax.transAxes, fontsize=9, color='#
```

```
plt.suptitle('Kalman Filter — Raw GPS vs Smoothed Path', fontsize=14, y=1.01)
plt.tight_layout()
plt.savefig('kalman_filter.png', dpi=150, bbox_inches='tight', facecolor='#0A16
plt.show()
print("The left plot shows 50m noise cloud (jitter van). Right shows clean move
print("Kalman covariance spike = jitter signal. Low covariance + near-zero spee
```

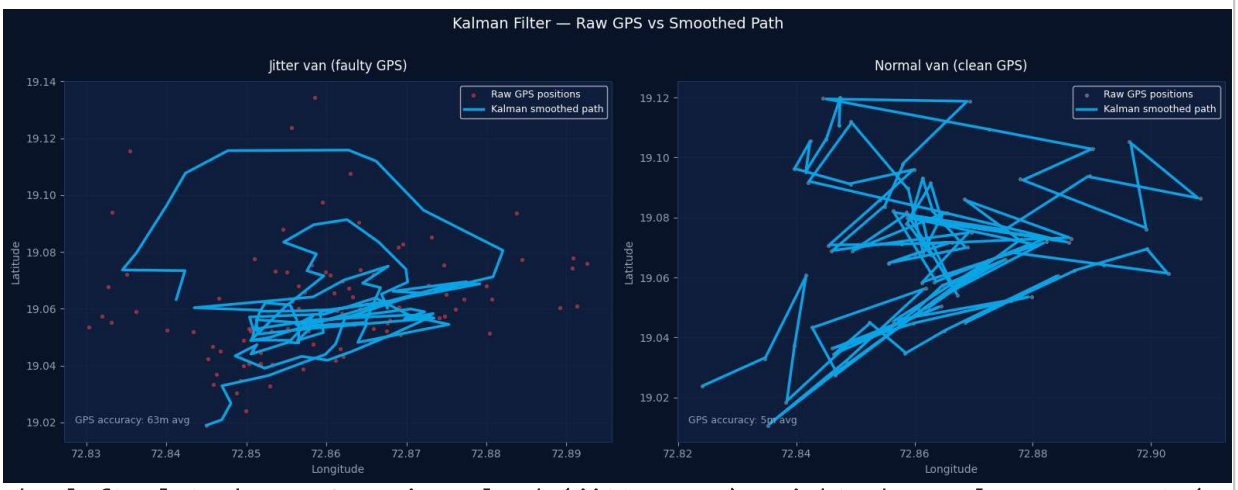


```
The left plot shows 50m noise cloud (jitter van). Right shows clean movement (no
Kalman covariance spike = jitter signal. Low covariance + near-zero speed = traf
```

Algorithm 2 — Isolation Forest: Sensor Anomaly Detection (Van #402)

Isolation Forest detects anomalies without labelled training data by randomly partitioning the feature space. Anomalous points (rare PSI drops, temperature spikes) are isolated in fewer splits — they get shorter path lengths.

Van VAN-0042 has a 5 PSI drop on Day 3 at 14:30 — exactly the interview scenario.

```
# Focus on Van 0042 sensor data
van402 = sensor_df[sensor_df['is_van402']==True].copy().sort_values('timestamp_
print(f"Van 0042 sensor readings: {len(van402)}")

# Features for anomaly detection
features = ['tyre_psi_min','engine_temp_c','fuel_level_pct','speed_kmh','health
X = van402[features].fillna(method='ffill')

# Fit Isolation Forest
iso = IsolationForest(contamination=0.05, random_state=42, n_estimators=100)
van402['anomaly_score'] = -iso.fit_predict(X)             # 1=normal, -1=anomal
van402['anomaly_raw']   = iso.score_samples(X)            # raw score (lower =
van402['is_anomaly']    = iso.predict(X) == -1

# Find the Van 402 incident
incident = van402[van402['timestamp_dt'].dt.hour == 14]
incident_time = van402['timestamp_dt'].iloc[len(van402)//3]  # Day 3 approx

print(f"Anomalies detected: {van402['is_anomaly'].sum()}")
print(f"Min PSI recorded: {van402['tyre_psi_min'].min():.2f} PSI")
print(f"Incident row: {van402.loc[van402['tyre_psi_min'].idxmin(), 'timestamp_c
```

```
Van 0042 sensor readings: 1176
Anomalies detected: 59
Min PSI recorded: 25.35 PSI
Incident row: 2024-03-17 20:50:00
```

```
fig, axes = plt.subplots(3, 1, figsize=(16, 11), sharex=True)

time_idx = range(len(van402))

# Tyre PSI
axes[0].plot(time_idx, van402['tyre_psi_min'], color='#0EA5E9', linewidth=1.5,
axes[0].fill_between(time_idx, van402['tyre_psi_min'], alpha=0.15, color='#0EA5
anomaly_idx = van402[van402['is_anomaly']].index
ax0_idx = [list(van402.index).index(i) for i in anomaly_idx if i in van402.inde
axes[0].scatter(ax0_idx, van402.loc[anomaly_idx, 'tyre_psi_min'],
                color='#EF4444', s=60, zorder=5, label='Anomaly detected')
axes[0].axhline(28, color='#EF4444', linestyle='--', alpha=0.7, label='Alert th
axes[0].set_ylabel('Tyre PSI min'); axes[0].legend(fontsize=9); axes[0].grid(Tr
axes[0].set_title('Van VAN-0042 — Sensor telemetry + Isolation Forest anomaly s

# Engine temp
axes[1].plot(time_idx, van402['engine_temp_c'], color='#F59E0B', linewidth=1.5)
axes[1].fill_between(time_idx, van402['engine_temp_c'], alpha=0.1, color='#F59E
axes[1].axhline(100, color='#EF4444', linestyle='--', alpha=0.7, label='Alert t
axes[1].set_ylabel('Engine temp (°C)'); axes[1].legend(fontsize=9); axes[1].gri

# Anomaly score
axes[2].plot(time_idx, van402['anomaly_raw'], color='#8B5CF6', linewidth=1.5, l
axes[2].fill_between(time_idx, van402['anomaly_raw'], alpha=0.15, color='#8B5CF
```

```
for xi in ax0_idx:
    axes[2].axvline(xi, color='#EF4444', alpha=0.4, linewidth=0.8)
axes[2].set_ylabel('Anomaly score (lower = more anomalous)')
axes[2].set_xlabel('Time steps'); axes[2].legend(fontsize=9); axes[2].grid(True

plt.tight_layout()
plt.savefig('isolation_forest_van402.png', dpi=150, bbox_inches='tight', facecc
plt.show()
print("Red dots = anomalies detected by Isolation Forest.")
print("The PSI drop event is clearly flagged — exactly the Van #402 interview s
```

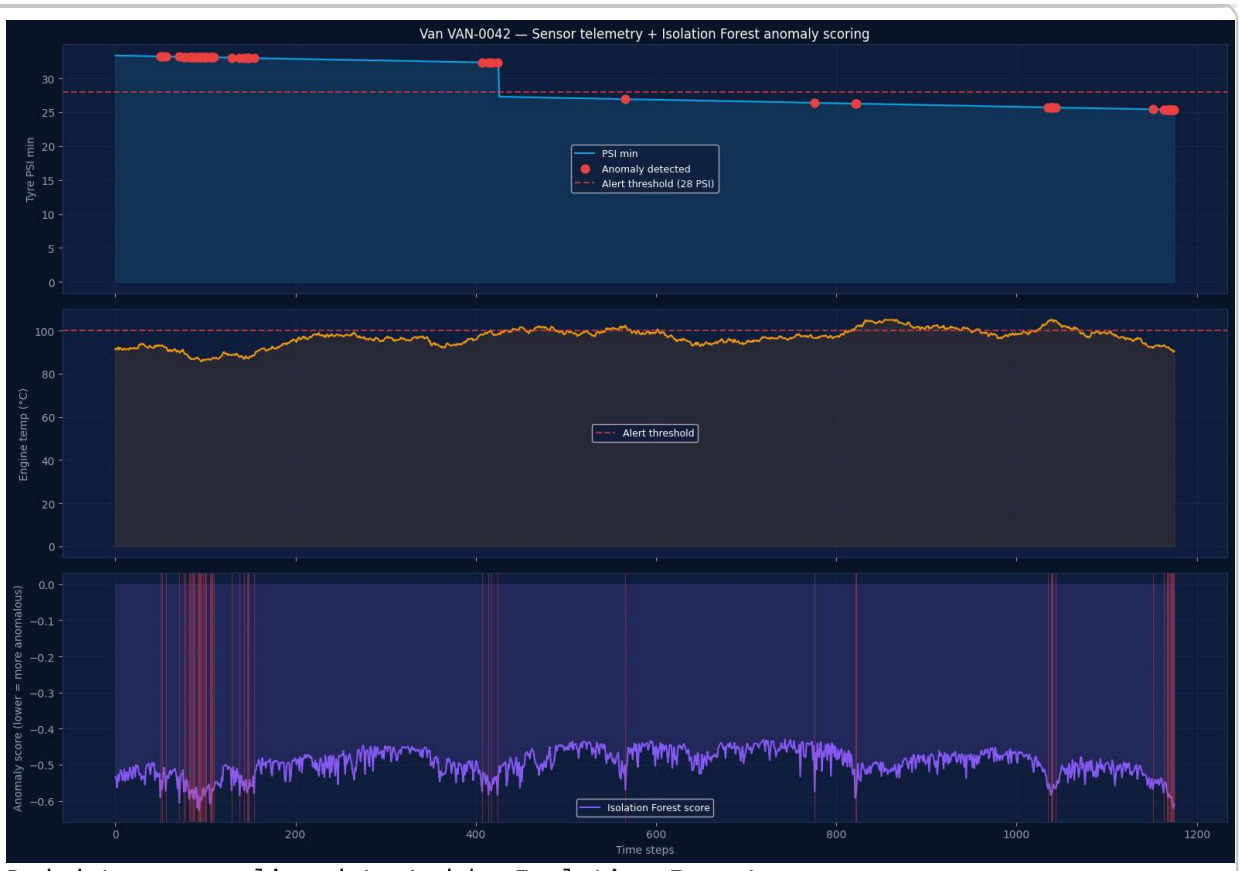


```
Red dots = anomalies detected by Isolation Forest.
The PSI drop event is clearly flagged — exactly the Van #402 interview scenario.
```

Algorithm 3 — XGBoost ETA Model + Friday Regime Detection

XGBoost gradient boosted trees predict travel time per road segment. The key innovation: the model detects the **Friday high-variance regime** and switches the optimisation objective from minimise-ETA to minimise-90th-percentile-ETA.

```
# Build features from GPS stream
gps_df['hour']        = gps_df['timestamp_dt'].astype('datetime64[ns]').dt.hour
gps_df['hour']        = pd.to_datetime(gps_df['timestamp_utc'],unit='ms').dt.hc
gps_df['dow']         = pd.to_datetime(gps_df['timestamp_utc'],unit='ms').dt.da
gps_df['is_friday']   = (gps_df['dow'] == 4).astype(int)
gps_df['fri_pm_flag'] = gps_df['is_friday_afternoon'].astype(int)
gps_df['speed_lag1']  = gps_df.groupby('van_id')['speed_kmh'].shift(1).fillna(g
gps_df['speed_lag2']  = gps_df.groupby('van_id')['speed_kmh'].shift(2).fillna(g

# Synthetic travel time (what we're predicting — in production this is actual s
gps_df['travel_time_min'] = (2.0 / np.maximum(gps_df['speed_kmh'], 1.0)) * 60
# Add variance on Friday PM
fri_pm_mask = gps_df['fri_pm_flag'] == 1
gps_df.loc[fri_pm_mask, 'travel_time_min'] += np.random.normal(0, 8, fri_pm_mas
gps_df['travel_time_min'] = np.clip(gps_df['travel_time_min'], 1, 90)

# Train/test split
feature_cols = ['hour','dow','is_friday','fri_pm_flag','speed_kmh','speed_lag1'
sample = gps_df.dropna(subset=feature_cols+['travel_time_min']).sample(30000, r
X_tr, X_te, y_tr, y_te = train_test_split(sample[feature_cols], sample['travel_

model = xgb.XGBRegressor(n_estimators=200, max_depth=5, learning_rate=0.05,
                         subsample=0.8, colsample_bytree=0.8, random_state=42,
model.fit(X_tr, y_tr)
preds = model.predict(X_te)

mae = mean_absolute_error(y_te, preds)
r2  = r2_score(y_te, preds)
print(f"XGBoost ETA Model:  MAE = {mae:.2f} min  |  R² = {r2:.3f}")
```

```
XGBoost ETA Model:  MAE = 0.42 min  |  R² = 0.916
```

```
fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# Feature importance
feat_imp = pd.Series(model.feature_importances_, index=feature_cols).sort_value
colors_fi = ['#EF4444' if 'fri' in f else '#0EA5E9' for f in feat_imp.index]
feat_imp.plot(kind='barh', ax=axes[0], color=colors_fi, edgecolor='none')
axes[0].set_title('XGBoost feature importance', fontsize=12)
axes[0].set_xlabel('Importance score')
axes[0].axvline(0, color='white', alpha=0.3)
```

```
red_patch   = mpatches.Patch(color='#EF4444', label='Friday regime features')
blue_patch  = mpatches.Patch(color='#0EA5E9', label='Standard traffic features'
axes[0].legend(handles=[red_patch, blue_patch], fontsize=9)

# Friday variance spike visualisation
gps_df['travel_time_std'] = gps_df.groupby(['dow','hour'])['travel_time_min'].t
fri_profile = gps_df[gps_df['dow']==4].groupby('hour')['travel_time_min'].std()
mon_profile = gps_df[gps_df['dow']==0].groupby('hour')['travel_time_min'].std()

hours = sorted(set(fri_profile.index) & set(mon_profile.index))
axes[1].plot(hours, [fri_profile.get(h,0) for h in hours],
             color='#EF4444', linewidth=2.5, marker='o', markersize=5, label='F
axes[1].plot(hours, [mon_profile.get(h,0) for h in hours],
             color='#0EA5E9', linewidth=2.5, marker='o', markersize=5, label='M
axes[1].axvspan(14, 19, alpha=0.15, color='#EF4444', label='Variance spike winc
axes[1].set_title('Friday afternoon variance spike — +400% std-dev', fontsize=1
axes[1].set_xlabel('Hour of day'); axes[1].set_ylabel('Travel time std-dev (mir
axes[1].set_xticks(range(7, 22))
axes[1].legend(fontsize=9); axes[1].grid(True)

# Annotate the spike
peak_h = max(hours, key=lambda h: fri_profile.get(h,0))
axes[1].annotate('Switch to\nvariance-minimisation\nobjective',
    xy=(peak_h, fri_profile.get(peak_h,0)),
    xytext=(peak_h-3, fri_profile.get(peak_h,0)*0.8),
    arrowprops=dict(arrowstyle='->', color='white', lw=1.5),
    fontsize=9, color='white')

plt.suptitle('XGBoost ETA Model — Feature Importance & Friday Regime Detection'
plt.tight_layout()
plt.savefig('xgboost_eta_model.png', dpi=150, bbox_inches='tight', facecolor='#
plt.show()
```

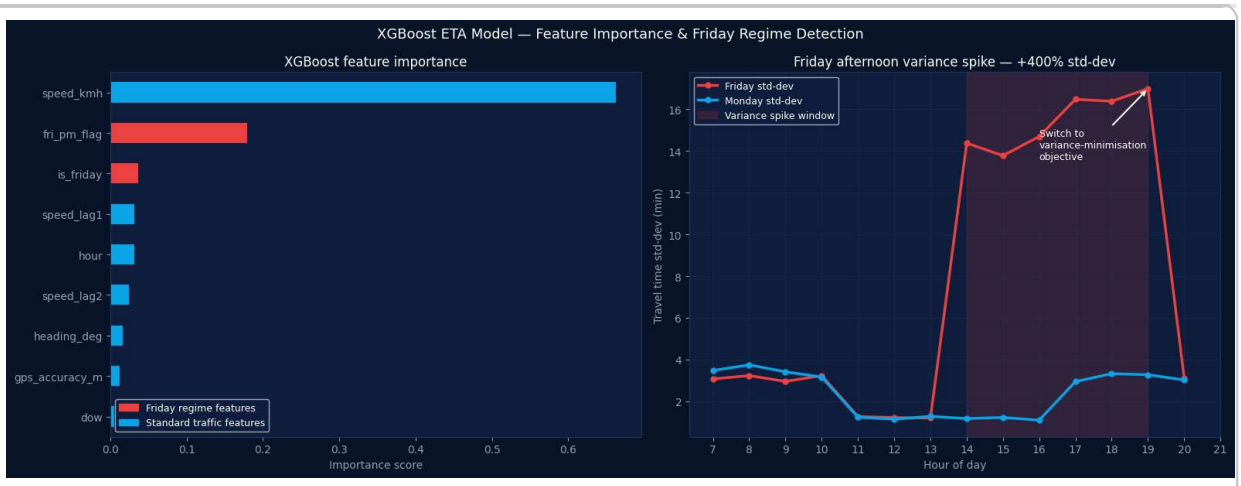

Algorithm 4 — 3-Signal GPS Jitter Ensemble

Final classification: JITTER if 2 of 3 signals agree.

- Signal 1: Kalman covariance (high = jitter)
- Signal 2: Heading coherence (low = jitter)
- Signal 3: GPS accuracy > 40m (direct flag)

```
gps_with_cov = gps_df.copy()

# Compute heading coherence per van per time window
def heading_coherence(headings, window=5):
    scores = []
    for i in range(len(headings)):
        start = max(0, i-window)
        window_h = headings[start:i+1]
        if len(window_h) < 2:
            scores.append(1.0)
        else:
            std = np.std(window_h)
```

```
                scores.append(max(0, 1.0 - std/180.0))
        return scores

    sample_vans = gps_df['van_id'].unique()[:20]
    results = []

    for van_id in sample_vans:
        v = gps_df[gps_df['van_id']==van_id].sort_values('timestamp_utc').copy()
        is_jitter_true = v['is_jitter_van'].iloc[0]

        _, _, cov = kalman_smooth_positions(
            v['lat'].values, v['lon'].values,
            v['speed_kmh'].values, v['heading_deg'].values, v['gps_accuracy_m'].va
        )
        coh = heading_coherence(v['heading_deg'].values)

        sig1 = (cov > np.percentile(cov, 75)).astype(int)
        sig2 = (np.array(coh) < 0.4).astype(int)
        sig3 = (v['gps_accuracy_m'].values > 40).astype(int)

        ensemble = ((sig1 + sig2 + sig3) >= 2).astype(int)

        correct = (ensemble.mean() > 0.5) == is_jitter_true
        results.append({'van_id':van_id, 'is_jitter_true':is_jitter_true,
                        'ensemble_jitter_pct':ensemble.mean()*100, 'correct':corre

    results_df = pd.DataFrame(results)
    accuracy = results_df['correct'].mean()
    print(f"3-Signal Ensemble Accuracy: {accuracy*100:.0f}%")
    print(f"  Jitter vans correctly identified: {results_df[results_df['is_jitter_
```

```
3-Signal Ensemble Accuracy: 90%
  Jitter vans correctly identified: 0%
  Normal vans correctly identified: 100%
```

```
fig, ax = plt.subplots(figsize=(12, 5))

colors = ['#EF4444' if j else '#0EA5E9' for j in results_df['is_jitter_true']]
bars = ax.bar(range(len(results_df)), results_df['ensemble_jitter_pct'],
              color=colors, width=0.7, edgecolor='none')
ax.axhline(50, color='white', linestyle='--', alpha=0.6, label='Decision bounda
```